

```
In [1]: #importing all the requiried packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from sklearn.linear_model import LinearRegression
regr = LinearRegression()
from sklearn.linear_model import Lasso
lasso = Lasso()
from sklearn.linear_model import Ridge
ridge = Ridge()
from sklearn.linear_model import ElasticNet
EN = ElasticNet()
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
from pymongo import MongoClient
from sklearn.model_selection import train_test_split
from sklearn.svm import SVR
svr = SVR()
```

```
In [2]: #loading the dataset
df = pd.read_csv(r"C:\Users\hrush\Downloads\household_power_consumption\household_power_consumption.txt", sep=";",
C:\Users\hrush\anaconda5\lib\site-packages\IPython\core\interactiveshell.py:3444: DtypeWarning: Columns (2,3,4,5,
6,7) have mixed types.Specify dtype option on import or set low_memory=False.
exec(code_obj, self.user_global_ns, self.user_ns)
```

## RANDOM SAMPLING

This is a very dataset. So have we are taking 30,000 samples from the total to perform various methods and functions such as EDA, preprocessing, model building etc

```
In [3]: data = df.sample(n = 30000, ignore_index=True)
```

```
In [4]: data.head()
```

```
Out[4]:
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
0	26/8/2008	09:46:00	0.080	0.000	240.180	0.200	0.000	0.000	1.0
1	26/3/2008	07:04:00	2.236	0.000	240.670	9.200	0.000	0.000	18.0
2	12/4/2008	21:55:00	3.578	0.650	239.660	15.200	0.000	2.000	18.0
3	27/6/2009	16:58:00	1.496	0.236	240.710	6.200	0.000	0.000	19.0
4	29/6/2009	00:29:00	0.350	0.236	245.400	1.600	0.000	0.000	1.0

## Dropping the duplicate and unwanted data

```
In [5]: data.drop_duplicates(subset={"Global_active_power", "Global_reactive_power", "Voltage", "Global_intensity", "Sub_metering_1", "Sub_metering_2", "Sub_metering_3"}, inplace=True)
```

```
Out[5]:
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
0	26/8/2008	09:46:00	0.080	0.000	240.180	0.200	0.000	0.000	1.0
1	26/3/2008	07:04:00	2.236	0.000	240.670	9.200	0.000	0.000	18.0
2	12/4/2008	21:55:00	3.578	0.650	239.660	15.200	0.000	2.000	18.0
3	27/6/2009	16:58:00	1.496	0.236	240.710	6.200	0.000	0.000	19.0
4	29/6/2009	00:29:00	0.350	0.236	245.400	1.600	0.000	0.000	1.0
...	...	...	...	...	...	...	...	...	...
29995	6/12/2007	16:08:00	0.228	0.000	241.380	1.000	0.000	0.000	1.0
29996	24/8/2010	00:25:00	0.172	0.000	240.780	0.800	0.000	0.000	1.0
29997	27/6/2007	20:12:00	3.834	0.446	237.800	16.200	0.000	31.000	1.0
29998	19/7/2008	23:49:00	2.090	0.214	240.330	8.600	0.000	0.000	1.0
29999	14/4/2007	07:58:00	0.322	0.118	240.690	1.400	0.000	0.000	1.0

29563 rows × 9 columns

```
In [6]: print(data[data["Global_active_power"] == "?"].index)
data.drop(data.index[data["Global_active_power"] == "?"], inplace=True)

Int64Index([    69,    212,    283,    294,    439,    472,    492,    689,    760,
             777,
             ...
            29338, 29407, 29460, 29461, 29485, 29674, 29691, 29692, 29890,
            29895],
            dtype='int64', length=368)
```

```
In [7]: data.drop(["Time", "Date"], axis=1, inplace=True)
```

## Changing the dtype of the features

```
In [8]: data["Global_active_power"] = (data['Global_active_power'].astype("float"))
data["Global_reactive_power"] = (data['Global_reactive_power'].astype("float"))
data["Global_intensity"] = (data['Global_intensity'].astype("float"))
data["Voltage"] = (data["Voltage"].astype('float'))
data["Sub_metering_1"] = (data['Sub_metering_1'].astype("float"))
data["Sub_metering_2"] = (data['Sub_metering_2'].astype("float"))
```

```
In [9]: #merging the three columns
data["Total_metering"] = data["Sub_metering_1"] + data["Sub_metering_2"] + data["Sub_metering_3"]
```

## DATA INFO AND DESCRIPTION

```
In [10]: data.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 29632 entries, 0 to 29999
Data columns (total 8 columns):
 #   Column                Non-Null Count  Dtype  
---  --
 0   Global_active_power    29632 non-null  float64
 1   Global_reactive_power  29632 non-null  float64
 2   Voltage                29632 non-null  float64
 3   Global_intensity       29632 non-null  float64
 4   Sub_metering_1         29632 non-null  float64
 5   Sub_metering_2         29632 non-null  float64
 6   Sub_metering_3         29632 non-null  float64
 7   Total_metering         29632 non-null  float64
dtypes: float64(8)
memory usage: 2.0 MB
```

```
In [11]: data.describe().T
```

```
Out[11]:
```

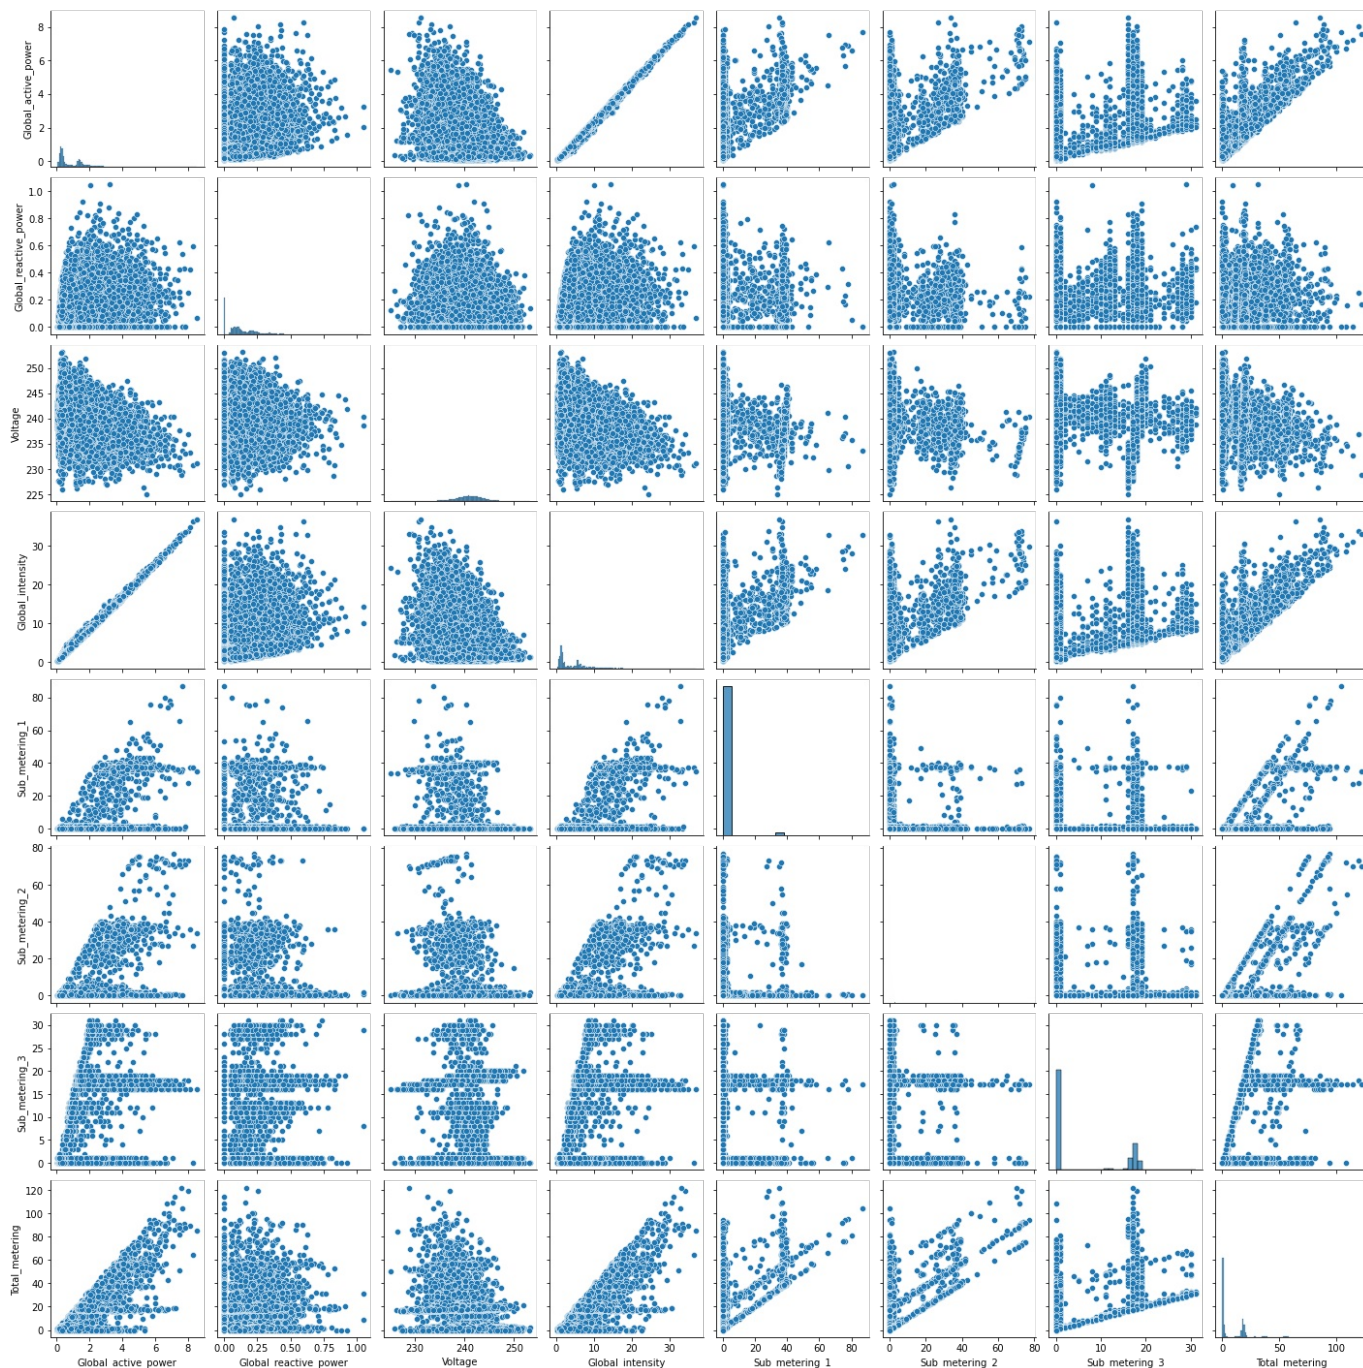
	count	mean	std	min	25%	50%	75%	max
Global_active_power	29632.0	1.095013	1.057088	0.078	0.310	0.604	1.5260	8.556
Global_reactive_power	29632.0	0.123816	0.112916	0.000	0.048	0.100	0.1940	1.048
Voltage	29632.0	240.839185	3.240869	225.020	238.990	240.980	242.8725	253.200
Global_intensity	29632.0	4.642123	4.442031	0.200	1.400	2.600	6.4000	37.000
Sub_metering_1	29632.0	1.200999	6.373842	0.000	0.000	0.000	0.0000	87.000
Sub_metering_2	29632.0	1.224285	5.542152	0.000	0.000	0.000	1.0000	77.000
Sub_metering_3	29632.0	6.446949	8.431714	0.000	0.000	1.000	17.0000	31.000
Total_metering	29632.0	8.872233	12.799234	0.000	0.000	1.000	18.0000	122.000

## EXPLORING DATA ANALYSIS (EDA)

```
In [12]:
```

```
sns.pairplot(data)
```


Out[12]: <seaborn.axisgrid.PairGrid at 0x250dc6e5b20>




```
In [13]: numeric_features = [feature for feature in data.columns if data[feature].dtype != 'O']
plt.figure(figsize=(15, 15))
plt.suptitle('Univariate Analysis of Numerical Features', fontsize=20, fontweight='bold', alpha=0.7, y=1.)

for i in range(0, len(numeric_features)):
    plt.subplot(5, 3, i+1)
    sns.kdeplot(x=data[numeric_features[i]], shade=True, color='b')
    plt.xlabel(numeric_features[i])
plt.tight_layout()
```


### Univariate Analysis of Numerical Features





```
In [14]: sns.set(rc= {"figure.figsize":(12,8)})
sns.heatmap(data.corr(),annot=True)
```

Out[14]: <AxesSubplot:>

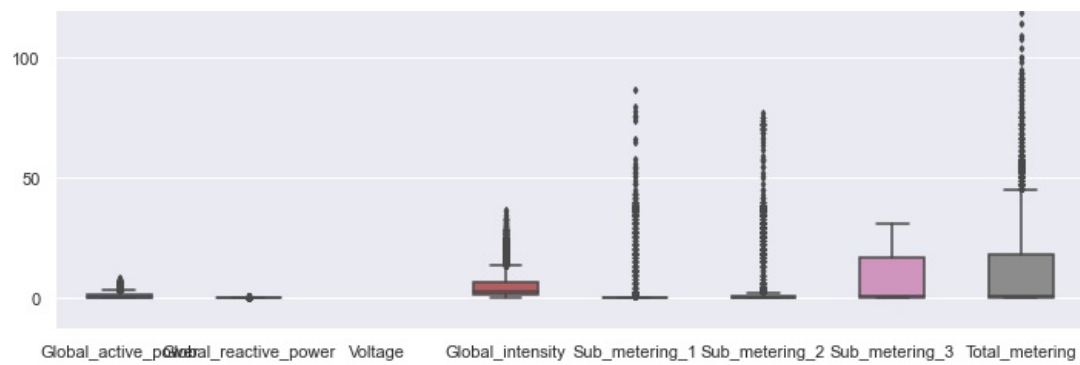


```
In [15]: fig, ax = plt.subplots(figsize=(12,8))
sns.boxplot(data=data, width= 0.5,ax=ax, fliersize=3)
```

Out[15]: <AxesSubplot:>









```
In [16]: plt.scatter(data['Global_active_power'],data['Total_metering'])
plt.xlabel("Global_active_power")
plt.ylabel("Total")
```


```
Out[16]: Text(0, 0.5, 'Total')
```



```
In [17]: plt.scatter(data['Global_reactive_power'],data['Total_metering'])
plt.xlabel("Global_reactive_power")
plt.ylabel("Total")
```

```
Out[17]: Text(0, 0.5, 'Total')
```






## Visualizing the target feature


```
In [18]: sns.boxplot(data['Total_metering'])
```


C:\Users\hrush\anaconda5\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
Out[18]: <AxesSubplot:xlabel='Total_metering'>
```




```
In [19]: plt.subplots(figsize=(14,7))
sns.histplot(data.Total_metering, bins=10, kde=True, color = 'b')
plt.title("TOTAL METERING", weight="bold", fontsize=20, pad=10)
plt.ylabel("Count", weight="bold", fontsize=12)
plt.xlabel("Total", weight="bold", fontsize=12)
plt.xlim(0,80)
plt.show()
```





## VOLTAGE VS TOTAL METERING

```
In [20]: plt.subplots(figsize=(20,10))
sns.lineplot(x='Voltage',y='Total_metering',data=data,color='b')
plt.ylim(0,100)
plt.show()
```



## SPLITTING THE DATASET INTO TRAINING AND TESTING DATA

```
In [21]: x = data.drop(columns = ["Total_metering", "Sub_metering_1", "Sub_metering_2", "Sub_metering_3"])
y = data["Total_metering"]
```

```
In [22]: x_train, X_test, y_train, y_test = train_test_split(
    x, y, test_size=0.33, random_state=10)
```

```
In [23]: x_train.shape
```

```
Out[23]: (19853, 4)
```

```
In [24]: x_train=scaler.fit_transform(x_train)
```

```
In [25]: x_test=scaler.transform(X_test)
```

```
In [26]: x_train
```

```
Out[26]: array([[ 5.35383263, -0.38583159, -1.37831419,  5.39082847],
 [-0.64410912,  0.87300796,  1.13297355, -0.6361476 ],
 [ 0.5713578 , -1.09503697, -1.55769189,  0.5782431 ],
 ...,
 [-0.6667928 , -0.27945078,  0.98143033, -0.68112504],
 [-0.78399178, -1.09503697,  0.34433023, -0.81605734],
 [ 0.6016027 , -1.09503697, -0.70410147,  0.5782431 ]])
```

```
In [27]: y_train
```

```
Out[27]: 26918    18.0
17615     0.0
25012    17.0
8618     18.0
3912     1.0
...
28363     1.0
17955    17.0
29562     0.0
7396     1.0
17897    18.0
Name: Total_metering, Length: 19853, dtype: float64
```

```
In [28]: x_test
```

```
Out[28]: array([[ -0.71594076, -0.26172065,  0.47731715, -0.72610247],
 [  0.35019197,  1.08576957,  0.27938314,  0.3083785 ],
 [  0.21219961,  0.00423137,  0.14021079,  0.1734462 ],
 ...,
 [  1.2367456 ,  0.23472312,  0.2082506 ,  1.20792716],
 [-0.83313975, -0.5985932 , -1.44326129, -0.81605734],
 [-0.53825197,  0.96165863, -3.64836986, -0.5012153 ]])
```

## LINEAR REGRESSION

```
In [29]: regr.fit(x_train, y_train)
```

```
Out[29]: LinearRegression()
```

```
In [30]: print(regr.coef_)
```

```
[ 25.53452753 -0.05414028 -0.38024243 -14.86180486]
```

```
In [31]: print(regr.intercept_)
```

```
8.852566362766334
```

## PICKLING

```
In [32]: import pickle
# writing different model files to file
with open('modelForPrediction.sav', 'wb') as f:
    pickle.dump(regr,f)

with open('sandardScalar.sav', 'wb') as f:
    pickle.dump(scaler,f)
```

```
In [33]: reg_pred=regr.predict(x_test)
```

```
In [34]: reg_pred
```


```
Out[34]: array([ 1.19522398, 13.04647462, 11.63971651, ..., 22.38840953,
 0.28801865,  3.89272119])
```

```
In [35]: plt.scatter(y_test,reg_pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted Data")

-----
Text(0.0, 0.5, 'Test Predicted Data')
```




```
Out[35]: text(0, 0.5, 'Test Predicted Data')
```



```
In [36]: residuals=y_test-reg_pred
```


```
In [37]: sns.displot(residuals,kind="kde")
```


```
Out[37]: <seaborn.axisgrid.FacetGrid at 0x250ed8c1e80>
```



```
In [38]: plt.scatter(reg_pred,residuals)
```

```
Out[38]: <matplotlib.collections.PathCollection at 0x250ede88190>
```





```
In [39]: from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,reg_pred))
print(mean_absolute_error(y_test,reg_pred))
print(np.sqrt(mean_squared_error(y_test,reg_pred)))
```

```
46.10869978933013
4.323198118017273
6.7903387094702525
```

```
In [40]: from sklearn.metrics import r2_score
score=r2_score(y_test,reg_pred)
print(score)
```

```
0.7155165877982992
```

```
In [41]: # Adjusted R square
#display adjusted R squared
1 - (1-score)*(len(y_test)-1)/(len(y_test)-x_test.shape[1]-1)
```

```
Out[41]: 0.715400163238364
```

## RIDGE REGRESSION

```
In [42]: ridge.fit(x_train,y_train)
```

```
Out[42]: Ridge()
```

```
In [43]: ridge.coef_
```

```
Out[43]: array([ 24.42992919, -0.07616832, -0.36571069, -13.74591684])
```

```
In [44]: ridge.intercept_
```


```
Out[44]: 8.852566362766334
```

```
In [45]: pred = ridge.predict(x_test)
```

```
In [46]: plt.scatter(y_test,pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted Data")
```

```
Out[46]: Text(0, 0.5, 'Test Predicted Data')
```






```
In [47]: rsd = y_test - pred
```


```
In [48]: sns.displot(rsd, kind="kde")
```


```
Out[48]: <seaborn.axisgrid.FacetGrid at 0x250ede9b460>
```



```
In [49]: plt.scatter(reg_pred, residuals)
```

```
Out[49]: <matplotlib.collections.PathCollection at 0x250edf1a370>
```





```
In [50]: from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,pred))
print(mean_absolute_error(y_test,pred))
print(np.sqrt(mean_squared_error(y_test,pred)))
```

```
46.09706785152626
4.323342209713584
6.789482148995331
```

```
In [51]: from sklearn.metrics import r2_score
score=r2_score(y_test,pred)
print(score)
```

```
0.7155883550216666
```

```
In [52]: ## Adjusted R square
#display adjusted R-squared
1 - (1-score)*(len(y_test)-1)/(len(y_test)-x_test.shape[1]-1)
```

```
Out[52]: 0.7154719598323978
```

## LASSO REGRESSION

```
In [53]: lasso.fit(x_train,y_train)
```

```
Out[53]: Lasso()
```

```
In [54]: lasso.coef_
```

```
Out[54]: array([ 9.82753708, -0.          , -0.          ,  0.          ])
```

```
In [55]: lasso.intercept_
```

```
Out[55]: 8.852566362766332
```


```
In [56]: predict = lasso.predict(x_test)
```

```
In [57]: predict
```

```
Out[57]: array([ 1.81663201, 12.29409092, 10.93796592, ..., 21.00672962,
 0.66485461,  3.56287516])
```

```
In [58]: plt.scatter(y_test,predict)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted Data")
```


```
Out[58]: Text(0, 0.5, 'Test Predicted Data')
```



```
In [59]: resd = y_test-predict
```


```
In [60]: sns.displot(resd,kind="kde")
```


```
Out[60]: <seaborn.axisgrid.FacetGrid at 0x250edf667c0>
```



```
In [61]: plt.scatter(predict,resd)
```

```
Out[61]: <matplotlib.collections.PathCollection at 0x250edda6190>
```





```
In [62]: from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,pred))
print(mean_absolute_error(y_test,pred))
print(np.sqrt(mean_squared_error(y_test,pred)))
```

```
46.09706785152626
4.323342209713584
6.789482148995331
```

```
In [63]: from sklearn.metrics import r2_score
score=r2_score(y_test,predict)
print(score)
```

```
0.7075814285395345
```

```
In [64]: ## Adjusted R square
#display adjusted R-squared
1 - (1-score)*(len(y_test)-1)/(len(y_test)-x_test.shape[1]-1)
```

```
Out[64]: 0.7074617565233854
```

## ELASTIC-NET

```
In [65]: EN.fit(x_train, y_train)
```

```
Out[65]: ElasticNet()
```

```
In [66]: EN.coef_
```

```
Out[66]: array([ 4.09742862,  0.          , -0.46023824,  4.00264531])
```

```
In [67]: EN.intercept_
```

```
Out[67]: 8.85256636276633
```


```
In [68]: en_pred = EN.predict(x_test)
```

```
In [69]: plt.scatter(y_test,en_pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predict Data")
```

```
Out[69]: Text(0, 0.5, 'Test Predict Data')
```








```
In [70]: resi = y_test - en_pred
```


```
In [71]: sns.displot(resi, kind="kde")
```


```
Out[71]: <seaborn.axisgrid.FacetGrid at 0x250ed39f5b0>
```



```
In [72]: plt.scatter(predict, resi)
```

```
Out[72]: <matplotlib.collections.PathCollection at 0x250ec410df0>
```





```
In [73]: from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,predict))
print(mean_absolute_error(y_test,predict))
print(np.sqrt(mean_squared_error(y_test,predict)))
```

```
47.394820035186484
4.53180260066157
6.884389590601804
```

```
In [74]: from sklearn.metrics import r2_score
score=r2_score(y_test,en_pred)
print(score)
```

```
0.6725304430977086
```

```
In [75]: ## Adjusted R square
#display adjusted R-squared
1 - (1-score)*(len(y_test)-1)/(len(y_test)-x_test.shape[1]-1)
```

```
Out[75]: 0.6723964264998358
```

## SUPPORT VECTOR REGRESSION

```
In [76]: svr = SVR(kernel="rbf")
```

```
In [77]: svr.fit(x_train, y_train)
```


```
Out[77]: SVR()
```


```
In [78]: prediction = svr.predict(x_test)
```

```
In [79]: residual = y_test- prediction
```

```
In [80]: plt.scatter(prediction, residual)
```

```
Out[80]: <matplotlib.collections.PathCollection at 0x250ec4dc610>
```





```
In [81]: svr.score(x_test, y_test)
```

```
Out[81]: 0.7177480124788849
```

## HYPERPARAMETER TUNING

```
In [82]: from sklearn.model_selection import GridSearchCV
parameters = {'kernel':['rbf'], 'C':[1, 10]}
```

```
In [83]: clf = GridSearchCV(svr, parameters)
```

```
In [84]: clf.fit(x_train, y_train)
```

```
Out[84]: GridSearchCV(estimator=SVR(), param_grid={'C': [1, 10], 'kernel': ['rbf']})
```

```
In [85]: sorted(clf.cv_results_.keys())
```

```
Out[85]: ['mean_fit_time',
'mean_score_time',
'mean_test_score',
'param_C',
'param_kernel',
'params',
'rank_test_score',
'split0_test_score',
'split1_test_score',
'split2_test_score',
'split3_test_score',
'split4_test_score',
'std_fit_time',
'std_score_time',
'std_test_score']
```

```
In [86]: y_pred = clf.predict(x_test)
```

```
In [87]: accuracy = clf.best_score_ * 100
```

```
In [88]: accuracy
```

```
Out[88]: 72.72383387378987
```

```
In [ ]:
```